

# LAB 1-DDPM

314552036 周子翔

## Task1

### I. Complete and Explain of all #TODO code

#### 1. Todo1

```
##### TODO #####
# DO (variable) layers: Any de this part. Sean, 22 hours ago • in
self.layers = nn.ModuleList()
prev_dim = dim_in
for hid_dim in dim_hids:
    self.layers.append(TimeLinear(prev_dim, hid_dim, num_timesteps))
    prev_dim = hid_dim
self.out_layer = nn.Linear(prev_dim, dim_out)
#####
```

- It stacks several TimeLinear layers, which are linear layers modulated by a time embedding (so the network can use timestep information).
- Each hidden layer's size is specified by dim\_hids.
- The final output layer maps the last hidden features to the output dimension (usually the same as the input, for noise prediction).

```
##### TODO #####
# DO NOT change the code outside this part.
h = x
for layer in self.layers:
    h = layer(h, t)
    h = F.silu(h)
out = self.out_layer(h)
#####
```

- The input  $\underline{x}$  (noisy data) is passed through each layer, with the current timestep  $\underline{t}$  provided to each TimeLinear layer.
- After each layer, a SiLU activation is applied for non-linearity.
- The final output is produced by the last linear layer, which predicts the noise for the given input and timestep.

#### 2. Todo2

```

##### TODO #####
# DO NOT change the code outside this part.
# Compute xt.
alphas_prod_t = extract(self.var_scheduler.alphas_cumprod, t, x0)
# DDPM公式:  $x_t = \sqrt{\alpha_{\bar{t}}} * x_0 + \sqrt{1 - \alpha_{\bar{t}}} * noise$ 
xt = torch.sqrt(alphas_prod_t) * x0 + torch.sqrt(1 - alphas_prod_t) * noise

#####

```

- For a given clean image  $x_0$  and timestep  $t$ , it adds noise according to the DDPM formula.
- $\text{alphas\_prod\_t}$  is the cumulative product of alphas up to timestep  $t$ .
- The output  $x_t$  is a noisy version of  $x_0$  at timestep  $t$ .

### 3. Todo 3

```

##### TODO #####
# DO NOT change the code outside this part.
# compute x_{t-1}.
if isinstance(t, int):
    t = torch.tensor([t]).to(self.device)

# 1. predict noise
eps_pred = self.network(xt, t)

# 2. Posterior mean
beta_t = extract(self.var_scheduler.betas, t, xt) #  $\beta_t$ 
alpha_t = extract(self.var_scheduler.alphas, t, xt) #  $\alpha_t = 1 - \beta_t$ 
alpha_bar_t = extract(self.var_scheduler.alphas_cumprod, t, xt) #  $\bar{\alpha}_t$ 
t_prev = (t - 1).clamp(min=0)
alpha_bar_t_prev = extract(self.var_scheduler.alphas_cumprod, t_prev, xt) #  $\bar{\alpha}_{t-1}$ 
mean = (
    torch.sqrt(alpha_bar_t_prev) * (xt - torch.sqrt(1 - alpha_bar_t) * eps_pred) / torch.sqrt(alpha_bar_t)
    + torch.sqrt(1 - alpha_bar_t_prev) * eps_pred
)

# 3. Posterior variance
var = beta_t * (1 - alpha_bar_t_prev) / (1 - alpha_bar_t)

# 4. Reverse step
if (t == 0).all():
    x_t_prev = mean
else:
    noise = torch.randn_like(xt)
    x_t_prev = mean + torch.sqrt(var) * noise

#####
return x_t_prev

```

- Predicts the noise in  $x_t$  using the network.
- Calculates the mean and variance for the posterior distribution at timestep  $t$ .
- Samples the previous timestep  $x_{t-1}$  using the mean and variance, adding noise unless at the final step.

### 4. Todo 4

```

##### TODO #####
# DO NOT change the code outside this part.
# sample x0 based on Algorithm 2 of DDPM paper.
xt = torch.randn(shape).to(self.device)
num_timesteps = self.var_scheduler.num_train_timesteps
for i in reversed(range(num_timesteps)):
    t = torch.full((xt.shape[0],), i, dtype=torch.long, device=self.device)
    xt = self.p_sample(xt, t)
x0_pred = xt
#####

```

- Starts from pure noise.
- Iteratively denoises through all timesteps using p\_sample.
- Returns the final denoised sample.

## 5. Todo 5

```

##### TODO #####
# DO NOT change the code outside this part.
# compute noise matching loss.
batch_size = x0.shape[0]

# 1) random choose timestep
t = (
    torch.randint(0, self.var_scheduler.num_train_timesteps, size=(batch_size,))
    .to(x0.device)
    .long()
)

# 2) get GT noise, and use q_sample to get x_t
eps = torch.randn_like(x0)
xt = self.q_sample(x0, t, noise=eps)

# 3) predict noise
eps_pred = self.network(xt, t)


# 4) MSE loss (eps, eps_pred)
loss = F.mse_loss(eps_pred, eps)

#####

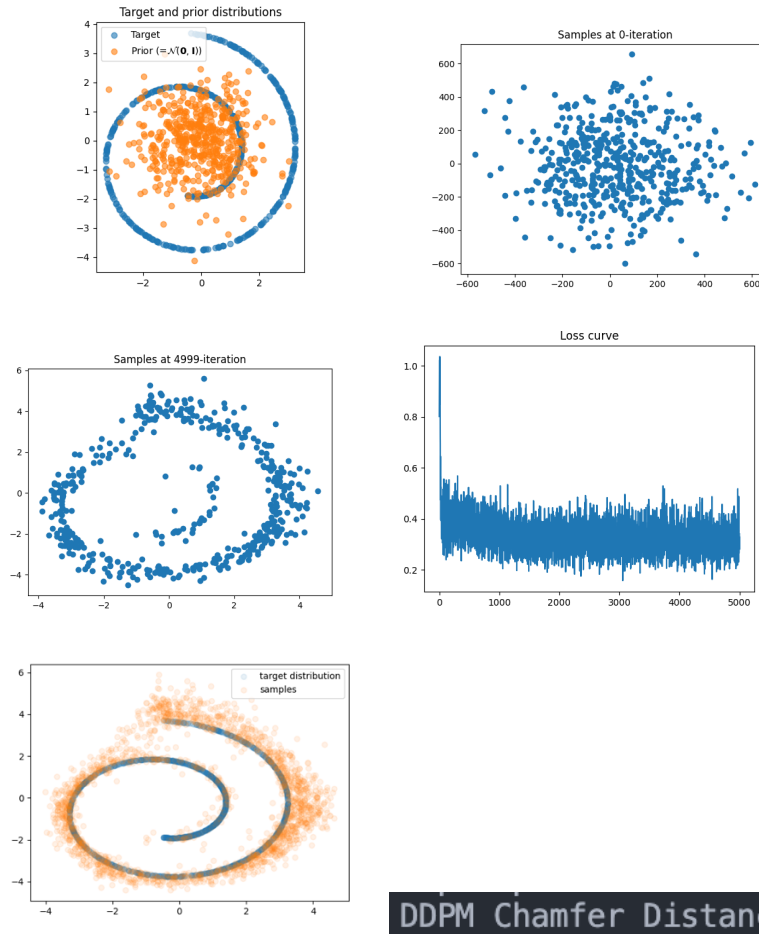
```

- Randomly selects a timestep for each sample in the batch.
- Adds noise to the clean data using q\_sample.
- Predicts the noise using the network.
- Computes the mean squared error between predicted and true noise.

## II. Fig of your implementation of q\_sample (10pts)



### III. Fig of loss curve and evaluation result (10 pts)



The DDPM implementation successfully learns the 2D swiss roll distribution, as shown by the sample plot and a reasonable Chamfer Distance.

### Task 2

#### IV. Complete and Explain of all #TODO code.

##### 1. Todo 1

```
##### TODO #####
# Assignment 1. Implement the DDPM forward step.
alpha_bar_t = extract(self.alphas_cumprod, t, x_0)
x_t = torch.sqrt(alpha_bar_t) * x_0 + torch.sqrt(1 - alpha_bar_t) * eps
#####
```

- Adds noise to the clean image according to the DDPM formula for a given timestep.

## 2. Todo 4

```
##### TODO #####
# Implement the cosine beta schedule (Nichol & Dhariwal, 2021).
# Hint:
# 1. Define  $\alpha_t = f(t/T)$  where  $f$  is a cosine schedule:
#      $\alpha_t = \cos^2( (t/T + s) / (1+s) ) * (\pi/2)$ 
#     with  $s = 0.008$  (a small constant for stability).
# 2. Convert  $\alpha_t$  into betas using:
#      $\beta_t = 1 - \alpha_t / \alpha_{t-1}$ 
# 3. Return betas as a tensor of shape [num_train_timesteps].
s = 0.008
steps = num_train_timesteps
t_arr = torch.arange(steps + 1, dtype=torch.float64)
f = lambda t: torch.cos(((t / steps + s) / (1 + s)) * np.pi / 2) ** 2
alphas_cumprod = f(t_arr)
betas = 1 - (alphas_cumprod[1:] / alphas_cumprod[:-1])
betas = torch.clip(betas, 0, 0.999)
betas = betas.to(torch.float32)
```

- It creates a smooth schedule for the cumulative product of alphas using a cosine function.
- Betas are then derived from the ratio of consecutive  $\alpha$  values.
- This schedule can improve sample quality compared to linear/quadratic schedules.

## 3. Todo 5

```
if predictor == "noise": #### TODO
    return self.step_predict_noise(x_t, t, net_out)
elif predictor == "x0": #### TODO
    return self.step_predict_x0(x_t, t, net_out)
elif predictor == "mean": #### TODO
    return self.step_predict_mean(x_t, t, net_out)
```

```

##### TODO #####
# 1. Extract beta_t, alpha_t, and alpha_bar_t from the scheduler.
if isinstance(t, int):
    t = torch.tensor([t], device=x_t.device)
beta_t = extract(self.betas, t, x_t)
alpha_t = extract(self.alphas, t, x_t)
alpha_bar_t = extract(self.alphas_cumprod, t, x_t)
alpha_bar_t_prev = extract(torch.cat([self.alphas_cumprod.new_ones(1), self.alphas_cumprod[:-1]], 0), t, x_t)
# 2. Compute the predicted mean  $\mu_{\theta}(x_t, t) = 1/\sqrt{\alpha_t} * (x_t - (\beta_t/\sqrt{1-\bar{\alpha}_t}) * \hat{\epsilon}_{\theta})$ .
mean = (1 / torch.sqrt(alpha_t)) * (x_t - (beta_t / torch.sqrt(1 - alpha_bar_t)) * eps_theta)
# 3. Compute the posterior variance  $\tilde{\beta}_t = ((1-\bar{\alpha}_{t-1})/(1-\bar{\alpha}_t)) * \beta_t$ .
posterior_var = ((1 - alpha_bar_t_prev) / (1 - alpha_bar_t)) * beta_t
# 4. Add Gaussian noise scaled by  $\sqrt{\tilde{\beta}_t}$  unless  $t == 0$ .
noise = torch.randn_like(x_t, device=x_t.device) if (t > 0).any() else torch.zeros_like(x_t, device=x_t.device)
mask = (t != 0).float().reshape(-1, *[1] * (x_t.dim() - 1)).to(x_t.device)
sample_prev = mean + mask * torch.sqrt(posterior_var) * noise
# 5. Return the final sample at t-1.
#####
return sample_prev

```

- Extracts scheduler parameters for the current timestep.
- Calculates the mean and variance for the reverse process.
- Adds noise unless at the final step.
- Returns the denoised sample.

```

##### TODO #####
if isinstance(t, int):
    t = torch.tensor([t], device=x_t.device)
beta_t = extract(self.betas, t, x_t)
alpha_t = extract(self.alphas, t, x_t)
alpha_bar_t = extract(self.alphas_cumprod, t, x_t)
alpha_bar_t_prev = extract(torch.cat([self.alphas_cumprod.new_ones(1), self.alphas_cumprod[:-1]], 0), t, x_t)
# Posterior mean (from DDPM paper)
mean = (
    torch.sqrt(alpha_bar_t_prev) * beta_t * x0_pred + torch.sqrt(alpha_t) * (1 - alpha_bar_t_prev) * x_t
) / (1 - alpha_bar_t)
posterior_var = ((1 - alpha_bar_t_prev) / (1 - alpha_bar_t)) * beta_t
noise = torch.randn_like(x_t, device=x_t.device) if (t > 0).any() else torch.zeros_like(x_t, device=x_t.device)
mask = (t != 0).float().reshape(-1, *[1] * (x_t.dim() - 1)).to(x_t.device)
sample_prev = mean + mask * torch.sqrt(posterior_var) * noise
#####
return sample_prev

```

- Uses the predicted clean image to compute the mean for the reverse process.
- Adds noise as in DDPM.

```


##### TODO #####
if isinstance(t, int):
    t = torch.tensor([t], device=x_t.device)
beta_t = extract(self.betas, t, x_t)
alpha_t = extract(self.alphas, t, x_t)
alpha_bar_t = extract(self.alphas_cumprod, t, x_t)
alpha_bar_t_prev = extract(torch.cat([self.alphas_cumprod.new_ones(1), self.alphas_cumprod[:-1]], 0), t, x_t)
posterior_var = ((1 - alpha_bar_t_prev) / (1 - alpha_bar_t)) * beta_t
noise = torch.randn_like(x_t, device=x_t.device) if (t > 0).any() else torch.zeros_like(x_t, device=x_t.device)
mask = (t != 0).float().reshape(-1, *[1] * (x_t.dim() - 1)).to(x_t.device)
sample_prev = mean_theta + mask * torch.sqrt(posterior_var) * noise
#####

```


- Implements the reverse step when the network directly predicts the posterior mean.

## V. Show the trajectory fig and compare results across the three beta schedulers.

1. Linear



2. Quadratic



3. Cosine




VI. Show the results of different predictors and discuss.

1. Noise

2.  $x_0$



3. posterior mean



VII. Screenshot of the Best FID of your training result, explain the training setting.  
The best FID score is 28.7 with linear schedulers and  $x_0$  predictors

FID: 28.72635241579301

- **Beta Schedule:** Linear (--mode linear)
- **Predictor:**  $x_0$  (--predictor  $x_0$ )
- **Batch Size:** 16 (--batch\_size 16)
- **Image Resolution:** 64x64 (--image\_resolution 64)
- **Number of Training Steps:** 50,000 (--train\_num\_steps 50000)
- **Diffusion Timesteps:** 1,000 (--num\_diffusion\_train\_timesteps 1000)
- **Beta Range:**  $\beta_1 = 1e-4$ ,  $\beta_T = 0.02$  (--beta\_1 1e-4 --beta\_T 0.02)
- **Seed:** 63 (--seed 63)
- **Optimizer:** Adam, learning rate 2e-4
- **Scheduler:** LambdaLR, warmup steps 200
- **Dataset:** AFHQ, max 3,000 images per category
- **Checkpoint/Results Directory:** Saved to Google Drive  
(/content/drive/MyDrive/results/predictor\_x0/beta\_linear/...)
- **Sampling Frequency:** Every 2,000 steps, model samples and saves checkpoint